

Cognitive Computer Vision

Computer Exercise

Tuesdays 14.15-15.45, Room D-118/119

Dr. Mikko Lauri

Email: lauri@informatik.uni-hamburg.de


SS 2018

Contents

1 Task 4 from Exercise 1

2 Exercise sheet 2

4. 2D Convolutions



- Sliding window (kernel) over image, calculate inner product with contents of input.

¹Image credit: <http://jeanvitor.com/convolution-parallel-algorithm-python/>

4. 2D Convolutions

```
1 tf.nn.conv2d(  
2     input ,  
3     filter ,  
4     strides ,  
5     padding ,  
6     use_cudnn_on_gpu=True ,  
7     data_format='NHWC' ,  
8     dilations=[1, 1, 1, 1] ,  
9     name=None)
```

- input – the input images: tensor of shape $[N \ H \ W \ C]$:
 N images, height H , width W , C channels.
- Example: one grayscale image, shape of input is
 $[1 \ H \ W \ 1]$

4. 2D Convolutions

```
1 tf.nn.conv2d(  
2     input ,  
3     filter ,  
4     ...)
```


- `filter` – the filtering kernels, tensor of shape $[F_h \ F_w \ N_i \ N_o]$: filter height F_h , width F_w , N_i input channels, N_o output channels.
- F_h and F_w typically odd, the middle element is the origin of the kernel.
- Example: one filter operating on single-channel input:
 $[F_h \ F_w \ 1 \ 1]$
- `filter[i,j,k,l]` is the weight of relative location (i,j) for the k th input channel to produce the l th output channel.

4. 2D Convolutions

```
1 tf.nn.conv2d(...,  
2   strides, ...)
```

- stride: number of units the filter shifts in each dimension, 1-D tensor with 4 elements: S_b, S_h, S_w, S_c

7 x 7 Input Volume



3 x 3 Output Volume



2

Figure: Stride 2 for height and width

²Image credit: <https://adeshpande3.github.io>

4. 2D Convolutions

```
1 tf.nn.conv2d(...,  
2   padding, ...)
```

- padding – How to pad image around the edges?
- "SAME": Pad the input with zeros around the edges to ensure output is of same shape as input.
- "VALID": no padding. Output shape will be smaller than input depending on the kernel shape.

4. 2D Convolutions

- `input` – the input images: tensor of shape $[N \ H \ W \ C]$: N images, height H , width W , C channels.
- `filter` – the filtering kernels, tensor of shape $[F_h \ F_w \ N_i \ N_o]$: filter height F_h , width F_w , N_i input channels, N_o output channels.
- `tf.reshape`, `tf.expand_dims`, or NumPy equivalents.

Static and dynamic shape of a tensor `x`

- The *static shape* is known at graph construction time and can be accessed by `x.shape`. Print it out anywhere!
- The *dynamic shape* can be partly or completely known only at runtime. Access by the operation `tf.shape(x)`. Use `sess.run!`
- The symbol `?` is shown for static shape where shape is unknown statically.

4. 2D Convolutions




Figure: Left: original, Right: Smoothed with $\sigma = 10$

4. 2D Convolutions

- You will need OpenCV + Python bindings to generate the filter kernel
- Use a classroom computer if you do not have OpenCV on your laptop
- If you have OpenCV on your laptop, to install Python bindings: `sudo apt-get install python-opencv`

1. Variables

- Values of tensors “live” only during a single `Session.run()` call, and are lost afterwards.
- Exception: `tf.Variable` stores data that persists over multiple runs
 - Model parameters
 - Bookkeeping: running averages and other statistics, number of iterations, ...
- Shape specified via `tf.initializer`; some variables are trainable (e.g., model parameters), others are not (e.g., number of iterations)

```
1 tf.Variable(initial_value=None,
2             trainable=True,
3             name=None,
4             ...)
```

2. Fully-connected layer

- Convenience function for creating a fully-connected (dense) neural network (NN) layer: `tf.layers.dense`

$$y = h(Wx + b)$$

- Input $x \in \mathbb{R}^m$, weights $W \in \mathbb{R}^{m \times k}$, bias $b \in \mathbb{R}^k$, output $y \in \mathbb{R}^k$, activation function $h: \mathbb{R} \rightarrow \mathbb{R}$
- Creates variables for W and b , adds ops to calculate y
- Infers the shape of W and b according to static shape of x and user-specified number of output neurons k

2. Fully-connected layer

- All layers in TF optimized to operate over several inputs at once (batch processing).
- `tf.layers.dense` expects an input of shape $[N \ m]$: number of inputs is N , each with dimensionality m – take this into account when defining the input!
- Same weight W and bias b will be applied to each of the N inputs.

3. Optimization and data logging

- Problem: minimize real-valued $F(\theta)$ where θ is a vector of parameters.
- Gradient descent idea:
 - Given the current parameters θ_n , calculate gradient of F at θ_n .
 - Gradient is a vector that points in the direction of the *steepest increase* of F : update θ_n by moving towards the direction of the *negative* gradient (direction of steepest *decrease*)
 - Repeat until convergence
- Gradient denoted by ∇F , step size $\gamma \in \mathbb{R}$, also called the *learning rate*.

$$\theta_{n+1} = \theta_n - \gamma \nabla F(\theta_n)$$

- There is no guarantee of finding a global optimum!

3. Optimization and data logging




Figure: Gradient at (x, y) illustrated by arrows originating from (x, y) .

3. Optimization and data logging

- Linear regression. Given *training data* (x_i, y_i) for $i = 1, \dots, N$,

$$\min_{a, b \in \mathbb{R}} \sum_{i=1}^N (ax_i + b - y_i)^2$$

- a and b make up our parameter vector θ . The expression above is the *loss function* to be minimized with respect to the parameters.

3. Optimization and data logging

- Tensorflow recipe for optimization:
 - 1 Choose a model architecture.
 - 2 Set up a computation graph to compute a prediction \hat{y}_i for given input x_i using the chosen architecture.
 - 3 Define a loss function that compares the prediction to the *ground truth* y_i : $L(\hat{y}_i, y_i)$
 - 4 Create an optimizer instance, e.g.,
`tf.train.GradientDescentOptimizer`
 - 5 Create a *training op* by calling the `minimize(loss)` method of the optimizer
 - 6 Iteratively run the training op while feeding training data.