

Cognitive Computer Vision

Computer Exercise

Tuesdays 14.15-15.45, Room D-118/119

Dr. Mikko Lauri

Email: lauri@informatik.uni-hamburg.de

SS 2018

Contents

1 Organization of exercises

Organization

- 6 computer exercises
- Course project: Build a visual saliency system (mandatory to pass course)



Figure: Left: Input image Right: Saliency map

Organization

- Exercises support and give you tools and methods to complete the course project.
- The exact instructions for the course project will be published soon. Deadline: June 12th
 - 1 You will work with a dataset and train a machine learning based visual saliency system.
 - 2 You will evaluate your results on a test dataset.
 - 3 You will present your system architecture and results to others (June 19th)
- Best saliency system will be awarded :)

Exercise topics

- Planned outline:
 - 1 Tensorflow introduction, 2D convolutions
 - 2 Variables and fully connected layers, training of neural networks
 - 3 Setting up a basic saliency system
 - 4 Set-up continued, evaluating a saliency system
 - 5 Evaluation continued
 - 6 Performance tuning

Tensorflow: Brief tutorial

- A tensor is a generalization of a matrix. Rank = number of dimensions

```
1 # rank 0 tensor; a scalar with shape []
2 3.
3 # rank 1 tensor; a vector with shape [3]
4 [1., 2., 3.]
5 # rank 2 tensor; a matrix with shape [2, 3]
6 [[1., 2., 3.], [4., 5., 6.]]
7 # rank 3 tensor with shape [2, 1, 3]
8 [[[1., 2., 3.]], [[7., 8., 9.]]]
```

Listing 1: Tensors

- Tensorflow is about setting up a graph of operations that take as input tensors, and output tensors.

Tensorflow graph

- A computational graph in TF consists of nodes and edges.
 - Nodes: operations
 - Edges: tensors (inputs and outputs of operations)
- Two phases:
 - 1 Set up a computational graph that defines what should be computed.
 - 2 Provide the required input values and run operations in the graph.
- When using TF, make sure you adhere to the phases: First set up everything in your graph. Only then run the required operations!

Tensorflow sessions

- The Python API of TF is merely used to construct the graph and specify which operations to run.
- Operations in the computational graph will actually run using the C++ backend of TF.
- A `tf.Session` allows communication back and forth with the backend, and running any operations.
- If a `tf.Graph` is like a `.py` file, a `tf.Session` is like the python executable.
- Now let's start the exercise to see these concepts in practice!

1. Getting started

- `tf.constant`: op that produces a constant value with given datatype as output.
- `+`: syntactic sugar for `tf.add` – op that takes two inputs and outputs elementwise sum of inputs. See also `*` and `/`
- `out = Session.run(fetches)` runs the op `fetches` (can also be a list or dictionary of ops) and returns the output.
Communication from TF backend to Python frontend.

Remember: Set up first, then run

First set up all ops in our computational graph, **only after that** start running ops, and don't add any more!

- Calling any `tf.operation` will add an operation to the computation graph. Operations can **not be removed** from the graph. Be very sure what you are doing if you don't follow this rule.


2. Visualizing the computation graph

- By default, unary functions will operate on tensors elementwise. Examples of unary functions:
`tf.sin`, `tf.cos`, `tf.square`, `tf.log`, ...
- `tf.matmul` is not `tf.multiply`: matrix vs. elementwise multiplication
- Good practice: give names to your ops

3. Feeding external data

- Recall: `out = Session.run(fetches)` runs the op `fetches` (can also be a list or dictionary of ops) and returns the output. Communication from TF backend to Python frontend.
- `out = Session.run(fetches, feed_dict)` allows user to give a dictionary of elements of type `placeholder`: value to feed to graph. Communication from Python frontend to TF backend.
- Example: `x=tf.placeholder(...),`
`feed_dict={x: [1.0, 2.0]}`

4. 2D Convolutions



- Sliding window (kernel) over image, calculate inner product with contents of input.

¹Image credit: <http://jeanvitor.com/convolution-parallel-algorithm-python/>

4. 2D Convolutions

```
1 tf.nn.conv2d(  
2     input ,  
3     filter ,  
4     strides ,  
5     padding ,  
6     use_cudnn_on_gpu=True ,  
7     data_format='NHWC' ,  
8     dilations=[1, 1, 1, 1] ,  
9     name=None)
```

- input – the input images: tensor of size $[N \ H \ W \ C]$: N images, height H , width W , C channels.
- Example: one grayscale image, size of input is $[1 \ H \ W \ 1]$

4. 2D Convolutions

```

1 tf.nn.conv2d(
2     input,
3     filter,
4     ...)

```


- `filter` – the filtering kernels, tensor of size $[F_h \ F_w \ N_i \ N_o]$: filter height F_h , width F_w , N_i input channels, N_o output channels.
- F_h and F_w typically odd, the middle element is the origin of the kernel.
- Example: one filter operating on single-channel input:
 $[F_h \ F_w \ 1 \ 1]$
- `filter[i,j,k,l]` is the weight of relative location (i,j) for the k th input channel to produce the l th output channel.

4. 2D Convolutions

```
1 tf.nn.conv2d(...,
2   strides, ...)
```

- stride: number of units the filter shifts in each dimension, 1-D tensor with 4 elements: S_b, S_h, S_w, S_c

7 x 7 Input Volume



3 x 3 Output Volume



2

Figure: Stride 2 for height and width

²Image credit: <https://adeshpande3.github.io>

4. 2D Convolutions

```
1 tf.nn.conv2d(...,  
2   padding, ...)
```

- padding – How to pad image around the edges?
- "SAME": Pad the input with zeros around the edges to ensure output is of same size as input.
- "VALID": no padding. Output size will be smaller than input depending on the kernel size.

4. 2D Convolutions

- Make sure your input and filter kernel are 4D tensors with appropriate size.
- `input` – the input images: tensor of size $[N \ H \ W \ C]$: N images, height H , width W , C channels.
- `filter` – the filtering kernels, tensor of size $[F_h \ F_w \ N_i \ N_o]$: filter height F_h , width F_w , N_i input channels, N_o output channels.
- `tf.reshape`, `tf.expand_dims`
- Output will also be 4D tensor